

The following model is the standard GAN which is part of **Exercise 1**. It is a very simple example and you can improve it by adding convolutions and many other ideas that we talked about if you want. Fill in the missing pieces and train it.

```
In [1]: %matplotlib inline

import os
import numpy as np
import math
import multiprocessing
import matplotlib.pyplot as plt

import torchvision.transforms as transforms
from torchvision.utils import save_image, make_grid
from torch.optim.optimizer import Optimizer, required
from torch.utils.data import DataLoader
from torchvision import datasets
from torch.autograd import Variable

import torch.nn as nn
import torch.nn.functional as F
import torch

device = "cuda" if torch.cuda.is_available() else "cpu"

os.makedirs("images_gan", exist_ok=True)
os.makedirs("images_cgan", exist_ok=True)

batch_size = 128                                     #size of the batches
lr = 0.0005                                           #adam: learning rate
b1 = 0.5                                              #adam: decay of first order moment
b2 = 0.999                                           #adam: decay of second order moment
n_cpu = multiprocessing.cpu_count()                  #number of cpu threads to use
latent_dim = 100                                     #dimensionality of the latent space
img_size = 28                                        #size of each image dimension
channels = 1                                          #number of image channels
sample_interval = 400                                #interval between image samples

img_shape = (channels, img_size, img_size)
torch.manual_seed(42)
```

Out[1]: <torch._C.Generator at 0x7fcff0ad5490>

Excercise 1. Generative Adversarial Networks (GANs) (34%)

Implement a GAN and train it on the MNIST dataset (a). Plot 10 samples generated using the GAN (b). Finally, train a classifier C that classifies MNIST images. Use that classifier to approximate the marginal distribution of the generator $p(y)$, (i.e. the probability that the GAN

samples a particular class). Visualize the distribution using a bar-plot (c). You can use one-to-one updates applying one gradient step for the discriminator and generator successively. If you want to keep this simple you can also just use linear layers.

```

In [2]: def to_onehot(digits, num_classes):
        """ [[3]] => [[0, 0, 1]]
        """
        labels_onehot = torch.zeros(digits.shape[0], num_classes).to(device)
        labels_onehot.scatter_(1, digits.view(-1, 1), 1)
        return labels_onehot

def plot_class_distributions(y_pred, num_classes):
    class_distributions = [np.sum(y_pred == i) for i in range(num_classes)]
    plt.bar(list(range(num_classes)), class_distributions, tick_label=list(range(num_classes)))
    plt.ylabel("Number of predictions")
    plt.xlabel("Class")
    plt.plot()

class GeneratorBlock(nn.Module):
    def __init__(self, in_feat, out_feat, activation, use_norm=True):
        super().__init__()
        self.dense = nn.Linear(in_feat, out_feat)
        self.activation = activation
        if use_norm:
            self.bn = nn.BatchNorm1d(out_feat, 0.8)
        else:
            self.bn = None

    def forward(self, x):
        x = self.dense(x)
        x = self.activation(x)
        if self.bn != None:
            x = self.bn(x)
        return x

class Generator(nn.Module):
    def __init__(self, num_classes=0):
        super().__init__()
        self.b1 = GeneratorBlock(latent_dim + num_classes, 128, activation=nn.LeakyReLU(0.2))
        self.b2 = GeneratorBlock(128, 256, activation=nn.LeakyReLU(0.2))
        self.b3 = GeneratorBlock(256, 512, activation=nn.LeakyReLU(0.2))
        self.b4 = GeneratorBlock(512, 1024, activation=nn.LeakyReLU(0.2))
        self.b5 = GeneratorBlock(1024, 1 * 28 * 28, activation=nn.Tanh())

    def forward(self, x, y=None):
        if y != None:
            x = torch.cat((x, y), dim=1)
        x = self.b1(x)
        x = self.b2(x)
        x = self.b3(x)
        x = self.b4(x)
        x = self.b5(x)
        x = x.view(-1, 1, 28, 28)
        return x

class Discriminator(nn.Module):
    def __init__(self, num_classes=0, use_sigmoid=True):

```

```
super().__init__()
self.dense_1 = nn.Linear(1 * 28 * 28 + num_classes, 512)
self.dense_2 = nn.Linear(512, 256)
self.dense_3 = nn.Linear(256, 1)
self.leaky_relu = nn.LeakyReLU(0.2)
self.use_sigmoid = use_sigmoid

def forward(self, x, y=None):
    x = x.view(x.shape[0], -1)
    if y != None:
        x = torch.cat((x, y), dim=1)
    x = self.dense_1(x)
    x = self.leaky_relu(x)
    x = self.dense_2(x)
    x = self.leaky_relu(x)
    x = self.dense_3(x)
    if self.use_sigmoid:
        x = torch.sigmoid(x)
    return x
```

```
In [3]: # Loss function
bce_loss = torch.nn.BCELoss()

# Initialize generator and discriminator
generator = Generator()
discriminator = Discriminator()

generator.to(device)
discriminator.to(device)
bce_loss.to(device)

# Configure data loader
os.makedirs("./mnist", exist_ok=True)
dataloader = torch.utils.data.DataLoader(
    torch.utils.data.ConcatDataset(
        [
            datasets.MNIST(
                "./mnist",
                train=True,
                download=True,
                transform=transforms.Compose(
                    [transforms.Resize(img_size), transforms.ToTensor()],
                ),
            ),
            datasets.MNIST(
                "./mnist",
                train=False,
                download=True,
                transform=transforms.Compose(
                    [transforms.Resize(img_size), transforms.ToTensor()],
                ),
            ),
        ]
    ),
    batch_size=batch_size,
    shuffle=True,
    num_workers=n_cpu
)

# Optimizers
optimizer_G = torch.optim.Adam(generator.parameters(), lr=lr, betas=(b1, b2))
optimizer_D = torch.optim.Adam(discriminator.parameters(), lr=lr, betas=(b1, b2))
```

```

In [4]: # -----
# Training
# -----

def train_loop(generator, discriminator, dataloader, optimizer_g, optimizer_d, num_epochs, iterations):
    gen_losses, disc_losses = [], []
    for epoch in range(num_epochs):
        running_gen_loss = 0.0
        running_disc_loss = 0.0
        for i, (real_imgs, y) in enumerate(dataloader):

            real_imgs = real_imgs.to(device)
            if with_labels:
                y = to_onehot(y.to(device), 10)
            else:
                y = None

            # -----
            # Train Generator
            # -----

            # WGANs train the discriminator more often than the generator
            if i % 5 == 0 or not use_wasserstein:
                optimizer_G.zero_grad()

                z = torch.randn((real_imgs.shape[0], latent_dim)).to(device)
                gen_imgs = generator(z, y)

                y_pred_fake = discriminator(gen_imgs, y)

                if use_wasserstein:
                    g_loss = loss_func(y_pred_fake)
                else:
                    g_loss = loss_func(y_pred_fake, torch.zeros_like(y_pred_fake))
                g_loss.backward()
                optimizer_G.step()

            # -----
            # Train Discriminator
            # -----

            optimizer_D.zero_grad()

            gen_imgs = generator(z, y)

            y_pred_real = discriminator(real_imgs, y)
            y_pred_fake = discriminator(gen_imgs, y)

            if use_wasserstein:
                d_loss = loss_func(y_pred_real) - loss_func(y_pred_fake)
            else:
                real_loss = loss_func(y_pred_real, torch.zeros_like(y_pred_real))
                fake_loss = loss_func(y_pred_fake, torch.ones_like(y_pred_fake))
                d_loss = (real_loss + fake_loss) / 2
            d_loss.backward()

```

```

optimizer_D.step()

# clip weights of discriminator when using WGAN
if use_wasserstein:
    for p in discriminator.parameters():
        p.data.clamp_(-0.01, 0.01)


running_gen_loss += g_loss.item()
running_disc_loss += d_loss.item()

batches_done = epoch * len(dataloader) + i
if batches_done % sample_interval == 0:
    # You can also save samples in your drive & maybe save y
    save_image(gen_imgs.data[:25], "images_gan/GAN-%d.png" %


gen_loss = running_gen_loss / iterations
disc_loss = running_disc_loss / iterations
loss = gen_loss + disc_loss
gen_losses.append(gen_loss)
disc_losses.append(disc_loss)
print(f"Epoch {epoch + 1}/{num_epochs} ==> loss: {loss}, gen_loss: {gen_loss}, disc_loss: {disc_loss}")
grid = make_grid(gen_imgs.data[:25], nrow=5, normalize=True).cpu_
# Channels first (PyTorch) to channels last (matplotlib)
grid = np.moveaxis(grid, 0, -1)
plt.imshow(grid, cmap='gray')
plt.axis('off')
plt.show()
return gen_losses, disc_losses

```

```
In [5]: gen_losses, disc_losses = train_loop(  
        generator,  
        discriminator,  
        dataloader,  
        optimizer_G,  
        optimizer_D,  
        bce_loss,  
        20,  
        False,  
        use_wasserstein=False  
    )
```



Epoch 20/20 ==> loss: 1.5051823876239698, gen_loss: 0.827243291154858, disc_loss: 0.677939096469112




```
In [7]: class Classifier(nn.Module):
    def __init__(self, in_dims):
        super().__init__()
        self.conv1 = nn.Conv2d(in_dims, 16, kernel_size=(3, 3), padding=1, stride=1)
        self.conv2 = nn.Conv2d(16, 32, kernel_size=(3, 3), padding=1, stride=1)
        self.conv3 = nn.Conv2d(32, 64, kernel_size=(3, 3), padding=1, stride=1)
        self.dense1 = nn.Linear(4 * 4 * 64, 10)

    def forward(self, x):
        x = self.conv1(x)
        x = torch.nn.functional.relu(x)
        x = self.conv2(x)
        x = torch.nn.functional.relu(x)
        x = self.conv3(x)
        x = torch.nn.functional.relu(x)
        x = x.view(-1, 4 * 4 * 64)
        x = self.dense1(x)
        x = torch.softmax(x, dim=1)
        return x

clf = Classifier(1).to(device)

clf_optim = torch.optim.Adam(clf.parameters(), lr=0.001)
criterion = torch.nn.CrossEntropyLoss().to(device)

for epoch in range(1, 11):
    running_loss = 0
    running_accuracy = 0
    iterations = 0
    for i, (x, y) in enumerate(dataloader):
        x, y = x.to(device), y.to(device)
        clf_optim.zero_grad()
        y_pred = clf(x)
        loss = criterion(y_pred, y)
        loss.backward()
        clf_optim.step()
        running_loss += loss.item()
        iterations += 1
    with torch.no_grad():
        accuracy = torch.mean(((torch.argmax(y_pred, 1) == y) * 1).float())
        running_accuracy += accuracy.item()
    loss = running_loss / iterations
    acc = running_accuracy / iterations
    print(f"Epoch {epoch}/10 ==> train loss: {loss}, train acc: {acc}")
```

```
Epoch 1/10 ==> train loss: 1.6623974343322532, train acc: 0.806054534
598286
Epoch 2/10 ==> train loss: 1.5195921829457257, train acc: 0.943880174
4520338
Epoch 3/10 ==> train loss: 1.4898731527642315, train acc: 0.972793973
6690556
Epoch 4/10 ==> train loss: 1.4841277424988406, train acc: 0.977984624
0005284
Epoch 5/10 ==> train loss: 1.4805172697063773, train acc: 0.981432815
```

```

35649
Epoch 6/10 ==> train loss: 1.4788157107407258, train acc: 0.983001844
5386312
Epoch 7/10 ==> train loss: 1.4771130475091323, train acc: 0.984458654
4106604
Epoch 8/10 ==> train loss: 1.4746645930916125, train acc: 0.987098867
3222348
Epoch 9/10 ==> train loss: 1.474323284473454, train acc: 0.9871315127
97075
Epoch 10/10 ==> train loss: 1.4732294851944694, train acc: 0.98815984
91773309

```


```

In [8]: z = torch.randn((1000, latent_dim)).to(device)

# Generate a batch of images
with torch.no_grad():
    gen_imgs = generator(z)
    y_pred = clf(gen_imgs)

y_pred_probs_ngan = y_pred
y_pred = np.argmax(y_pred.cpu().numpy(), axis=1)
plot_class_distributions(y_pred, 10)

```



Excercise 2. Wasserstein GANs (46%)

Unless you have been lucky, you should have seen a mode collapse (one class being sampled much more often than another). Now Wasserstein GANs are not only proven to generate better images, but they also do not have such a bad mode collapse. However, we need to enforce 1-lipschitz continuity of the discriminator D in order for WGANs to work. You can either implement Spectral-Norm-Layers¹ or do Gradient Penalty² to ensure this. Train your model on the data set. (33%)

```
In [9]: def wasserstein_loss(y_pred):
        return torch.mean(y_pred)

        # It's recommended to use RMSProp for WGAN
        generator = Generator()
        discriminator = Discriminator(0, use_sigmoid=False)

        optimizer_G = torch.optim.RMSprop(generator.parameters(), lr=0.00005)
        optimizer_D = torch.optim.RMSprop(discriminator.parameters(), lr=0.00005)

        generator.to(device)
        discriminator.to(device)
```

```
Out[9]: Discriminator(
  (dense_1): Linear(in_features=784, out_features=512, bias=True)
  (dense_2): Linear(in_features=512, out_features=256, bias=True)
  (dense_3): Linear(in_features=256, out_features=1, bias=True)
  (leaky_relu): LeakyReLU(negative_slope=0.2)
)
```

```
In [10]: gen_losses, disc_losses = train_loop(
        generator,
        discriminator,
        dataloader,
        optimizer_G,
        optimizer_D,
        wasserstein_loss,
        200,
        False,
        use_wasserstein=True
    )
```



Epoch 181/200 ==> loss: -0.830482601901513, gen_loss: -0.6982108256498682, disc_loss: -0.1322717762516447




**(a) Calculate a marginal distribution, again, and create a bar-plot.
(13%)**

```
In [11]: z = torch.randn((1000, latent_dim)).to(device)

# Generate a batch of images
with torch.no_grad():
    gen_imgs = generator(z)
    y_pred = clf(gen_imgs)

y_pred_probs_wgan = y_pred
y_pred = np.argmax(y_pred.cpu().numpy(), axis=1)
plot_class_distributions(y_pred, 10)
```



(b) Calculate the Inception Score (IS) for both models. (10%)

```

In [17]: def inception_score(y_pred, chunk_size=10):
    inc_scores = []
    for chunk_start in range(0, y_pred.shape[0], chunk_size):
        y_pred_cur = y_pred[chunk_start: chunk_start + chunk_size]
        # calculate marginal probability, which is just the average prob
        marginal_dist = torch.mean(y_pred_cur, dim=0)

        # calculate kl divergence for each sample between it's probability and marginal_dist
        # Add a small eps in case a class has probability 0.00, because log(0) is -inf
        eps = 1e-14
        kl_div = y_pred_cur * (torch.log(y_pred_cur + eps) - torch.log(marginal_dist + eps))

        # sum over all classes
        kl_div_samples = torch.sum(kl_div, dim=0)

        # average over all samples
        avg_kl = torch.mean(kl_div_samples)
        inc_score = torch.exp(avg_kl)
        inc_scores.append(inc_score.cpu().numpy())
    return np.mean(inc_scores)

inc_score_ngan = inception_score(y_pred_probs_ngan)
inc_score_wgan = inception_score(y_pred_probs_wgan)
print(f"GAN has an inception score of {inc_score_ngan}")
print(f"WGAN has an inception score of {inc_score_wgan}")

```

```

GAN has an inception score of  5.711552143096924
WGAN has an inception score of  5.353795051574707

```

Excercise 3. Conditional GANs (CGANs) (20%)

Now if we want to ultimately prevent mode collapse, we need to provide the information of what to sample additionally. This is what CGANs do. Implement a CGAN and train it on the data set. (15%)

```
In [13]: # Initialize generator and discriminator
generator = Generator(10)
discriminator = Discriminator(10)

generator.to(device)
discriminator.to(device)

# Optimizers
optimizer_G = torch.optim.Adam(generator.parameters(), lr=lr, betas=(b1, b2))
optimizer_D = torch.optim.Adam(discriminator.parameters(), lr=lr, betas=(b1, b2))

gen_losses, disc_losses = train_loop(
    generator,
    discriminator,
    dataloader,
    optimizer_G,
    optimizer_D,
    bce_loss,
    20,
    True,
    use_wasserstein=False
)
```



Epoch 20/20 ==> loss: 1.4606071692062252, gen_loss: 0.7738216651205171, disc_loss: 0.6867855040857082



```


In [19]: z = torch.randn((1000, latent_dim)).to(device)

y = [i % 100 for i in range(1000)]

# Generate a batch of images
with torch.no_grad():
    y = torch.LongTensor([i % 10 for i in range(1000)]).to(device)
    y = to_onehot(y, 10)
    gen_imgs = generator(z, y)
    y_pred = clf(gen_imgs)

y_pred_probs = y_pred
y_pred = np.argmax(y_pred.cpu().numpy(), axis=1)
plot_class_distributions(y_pred, 10)

```



```

In [21]: inc_score_cgan = inception_score(y_pred_probs)
print(f"GAN has an inception score of {inc_score_ngan}")
print(f"WGAN has an inception score of {inc_score_wgan}")
print(f"CGAN has an inception score of {inc_score_cgan}")

```

```

GAN has an inception score of  5.711552143096924
WGAN has an inception score of  5.353795051574707
CGAN has an inception score of  9.517807006835938

```

As we see the CGAN is clearly our best model.

(a) Sample a few pictures of the classes of your choice. (5%)

```
In [41]: fig, ax = plt.subplots(nrows=2, ncols=5, figsize=(25, 10))

for label in range(10):
    z = torch.randn((25, latent_dim)).cuda()
    with torch.no_grad():
        y = torch.ones((25,)).long().cuda() * label
        y = to_onehot(y, 10)
        gen_imgs = generator(z, y)
    grid = make_grid(gen_imgs, nrow=5, normalize=True).cpu().numpy()
    # Channels first (PyTorch) to channels last (matplotlib)
    grid = np.moveaxis(grid, 0, -1)
    ax[label // 5][label % 5].imshow(grid, cmap='gray')
    ax[label // 5][label % 5].axis('off')
    ax[label // 5][label % 5].set_title(label)
plt.show()
```

